

Test Plans — OrGanuz AI Site Automation

This folder holds a written test plan for every automated test group in the repository. Each plan describes four things:

- **Scope** — what the group covers.
- **Project / target / environment** — the Playwright project it runs as, and the site or backend it runs against.
- **Gating** — the specific conditions under which a test is allowed to skip.
- **Cases** — a table of every test in the group and what each one checks.

The grouping matches the per-test-group skills in `.claude/skills/` and the "Test Groups → Skills" map in `Architecture.html`. Whenever a test is added or removed, keep these plans in sync with the specs and with the counts in `CLAUDE.md`, following the `test-suite-parity` skill.

Groups

#	PLAN	PROJECT(S)	SPECS	CASES	TARGET
1	UI Sanity	chromium (disabled)	tests/ui/**	12	Marketing site (prod www.organuz.ai)
2	Product Public Sanity	product	tests/product/api/**	8	Dev calculator (no login)
3	Product Matrix & Role Contracts	product	tests/product/matrix/**	23	Offline — checked- in fixtures
4	Product Roles E2E	product-setup → product- authenticated (disabled)	tests/product/flows/**	13	Dev calculator (per- role login)
5	Organuz API Contracts	organuz-api	tests/organuz-api/**	1	Organuz Supabase / PostgREST
6	QA Agent	agent	tests/agent/**	2	Offline — stubs, no network/browser
7	External API Monitoring	monitoring (opt-in)	tests/monitoring/**	50	Live third-party maps — Govmap + Ofek

#	PLAN	PROJECT(S)	SPECS	CASES	TARGET
8	Security (penetration testing)	security	tests/security/**	30	Organuz Supabase / PostgREST (anon key)
9	Local Web (local-only marketing e2e)	local-web	tests/local-web/**	50	Marketing site (prod <code>www.organuz.ai</code>) — self-skips on CI
10	Accessibility	accessibility	tests/accessibility/**	30	Marketing site (prod <code>www.organuz.ai</code>) — local + CI
11	Product Fraud / ATO	fraud	tests/fraud/**	14	Product app origin + auth backend (dev by default)
12	Product Interactive UI	product	tests/product/auth/**, mobile/**, wizard/**	24	Dev calculator — cellular login, mobile, and wizard

Plans **1 and 4 are currently DISABLED** — their projects are commented out in `playwright.config.ts` (spec files retained; re-enable by uncommenting). Their case counts above describe the specs as written, not the current default run.

Default run (`npx playwright test`): **188 discovered tests** across seven projects — `product` 61 (public sanity 8 + matrix/role contracts 23 + English 6

- interactive UI 24), `local-web` 50, `accessibility` 30, `security` 30, `fraud` 14, `agent` 2, and `organuz-api` 1. Plans **1 and 4 are disabled** (their projects are commented out in `playwright.config.ts`), so the marketing UI sanity and the product roles e2e do not run in the default suite. Re-enable a group by uncommenting its project; the plan for each group documents which project to uncomment and the specs it covers.

The `local-web` group (Plan 9) is in the default local run but **self-skips on CI** — every spec `test.skip`s when `process.env.CI` is set — so those 50 specs run locally and stay green (skipped) on GitHub Actions. This is a sanctioned local/CI divergence; see the `test-suite-parity` skill.

Opt-in monitoring (`npm run test:monitoring`, i.e. `MONITORING_ENABLED=true`): this adds the `monitoring` project — 50 live availability checks against Govmap and Ofek (govmap 25 + ofek 25). It is **never** part of the default suite, so an outage at Govmap or Ofek can't break the PR gate. Instead it runs on the scheduled **External API Monitoring** workflow and alerts through an auto-managed GitHub issue plus Slack. With monitoring enabled, the total is **238** (default 188 + monitoring 50).

Sanctioned skips (never failures)

- **Group 7 (external API monitoring):** a spec skips when the govmap.gov.il edge serves an HTML block/challenge page to the runner (a geo/bot block — e.g. a CI runner outside Israel). A per-worker canary in `tests/monitoring/support/availability.ts` detects it and skips via each `beforeEach`. A **real** break — a 5xx, a connection error, or a real-but-wrong (non-HTML) payload — still fails; that failure is the alert.
- **Group 2 (public sanity):** skips only during a genuine dev outage (dev gateway or password gate down, no UI token extractable). A token that *is* observed but comes back malformed is a real regression and fails.
- **Group 4 (roles e2e) — disabled:** each role skips on its own — no `<ROLE>_PHONE` / `_OTP_CODE` credential (product-setup) or no saved session (product-authenticated).
- **Group 9 (local web):** every spec self-skips when `process.env.CI` is set — a top-level `localOnly()` guard `test.skip`s the whole file on CI. The 50 specs run (and pass) on a developer machine and are simply skipped on GitHub Actions; the `local-web` project is not in the CI matrix. Sanctioned local/CI divergence.
- **Group 12 (product interactive UI):** live geocode-driving wizard cases self-skip on CI; the real OTP-send case requires `PRODUCT_OTP_UI=true`; authenticated characterization requires `PRODUCT_WIZARD_E2E=true`; and the temporary mobile navigation diagnostic requires `MOBILE_NAV_EXPLORE=true`.

Everything else must pass. See `CLAUDE.md` → "Conventions" for the full policy.

PDFs

A PDF of every plan (and of this index) lives in `pdf/`. Regenerate them after editing any plan:

```
npm run test-plans:pdf # renders docs/test-plans/*.md → docs/test-plans/pdf/*.pdf
```

The conversion runs in two steps: Markdown → HTML via `marked`, then HTML → PDF via Playwright's headless Chromium. No external tool is required.